

Advanced SQL and Connector /J

Amos Azaria

Variables

- User-defined variables (@):
 - SET @passGrade = 60
 - SELECT @avgGrade := AVG(grade) FROM grades
 - SELECT AVG(grade) INTO @avgGrade FROM grades
- Local variables in stored procedures (will learn later)
 - DECLARE passGrade INT
- To see a variable's value you can SELECT it:
 - SELECT @avgGrade

Note the syntax! (in SQL Server it works with just =)

@avgGrade
70.0000000000

TEMPORARY TABLE

- Variables cannot hold tables.
- If you would like to use a table during execution, you can use the TEMPORARY keyword:
 - `CREATE TEMPORARY TABLE tempTable (id INT, name VARCHAR(1000));`
 - `INSERT INTO tempTable (SELECT id, lastName FROM students);`
 - (5 rows effected)
- You can also combine create with select:
 - `CREATE TEMPORARY TABLE tempTable2 AS (SELECT * FROM students);`

Aliases

- `SELECT * FROM students AS s INNER JOIN grades AS g ON s.id=g.studentId;`
- `SELECT * FROM students s INNER JOIN grades g ON s.id=g.studentId;`

	id	age	gender	degree	firstName	lastName	avg grade	courseId	studentId	grade	passed
• S	111	21	1	1	Chaya	Glass	73.3333333333	20	111	43	0
	111	21	1	1	Chaya	Glass	73.3333333333	30	111	90	1
S	111	21	1	1	Chaya	Glass	73.3333333333	50	111	87	1
	222	28	1	3	Tal	Negev	HULL	20	222	85	1
S	222	28	1	3	Tal	Negev	HULL	40	222	72	1
	333	24	0	1	Gadi	Golan	HULL	40	333	45	0
	444	23	0	1	Moti	Cohen	HULL	30	444	95	1

111	21	50	111	87	1
222	28	20	222	85	1
222	28	40	222	72	1
333	24	40	333	45	0
444	23	30	444	95	1

Aliases (cont.)

- `SELECT age+10 FROM students;`
- `SELECT age+10 AS future_age FROM students;`

age+10	future age
31	31
38	38
34	34
33	33
36	36



- `SELECT studentId, MAX(av) FROM (SELECT studentId, AVG(grade) AS av FROM grades GROUP BY studentId) AS t;`

✖ 4 10:25:50 SELECT studentId, MAX(av) FROM (SELECT studentId, AVG(grade) AS av FROM grades GROUP BY studentId) AS t; Error Code: 1248. Every derived table must have its own alias

Error Code: 1248. Every derived table must have its own alias

studentId	MAX(av)
111	95.0000

Transactions

- Suppose we want to transfer money from one bank account to another:
 - SET @transferAmount = 1000;
 - SELECT @firstBalance = amount FROM bankBalances WHERE userId = 777;

 - UPDATE bankBalances SET amount = @ firstBalance - @ transferAmount WHERE userId = 777;
 - SELECT @secondBalance = amount FROM bankBalances WHERE userId = 888;

 - UPDATE bankBalances SET amount = @secondBalance + @ transferAmount WHERE userId = 888;
- What might be the problem with this execution?

Transactions (cont.)

- Transactions are guaranteed to be executed completely or nothing at all (**A**tomicity in ACID). In this example we also rely on the Isolation attribute (and **D**urability).
- When using a single query, it is treated as a transaction.
- We can combine several queries into a single transaction by using `START TRANSACTION` and ending with `COMMIT`.

Transactions (cont.)

```
SET @transferAmount = 1000;
```

```
START TRANSACTION
```

```
SELECT @firstBalance = amount FROM bankBalances  
WHERE userId = 777;
```

```
UPDATE bankBalances SET amount = @firstBalance - @  
transferAmount WHERE userId = 777;
```

```
SELECT @secondBalance = amount FROM bankBalances  
WHERE userId = 888;
```

```
UPDATE bankBalances SET amount = @secondBalance +  
@transferAmount WHERE userId = 888;
```

```
COMMIT
```

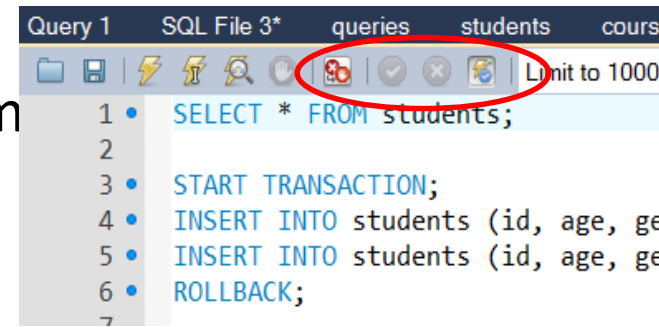

Transactions (ROLLBACK)

- If you have any error during the transaction, you can call ROLLBACK to undo the current transaction (until previous COMMIT).

Transactions in WorkBench

- There are several transaction related buttons on the workbench toolbar:

- Continue even if an error occurs.
- Autocommit every query. (set autocom



Show ROLLBACK /
COMMIT example in
WorkBench

- I found it very difficult to write a transaction in workbench with a ROLLBACK in case of error (and COMMIT otherwise). Though, this is the regular behavior if using the MySQL command-line client, and we later see an example in a stored procedure.

Stored Procedures

- Procedures that are stored inside the database:
 - Can be accessed from different programming languages.
 - Can save communication time
 - Can be modified 'on the fly' (changes can be made without restarting the server)

Stored Procedure - Example

```
DELIMITER $$
```

```
CREATE PROCEDURE student_avg
```

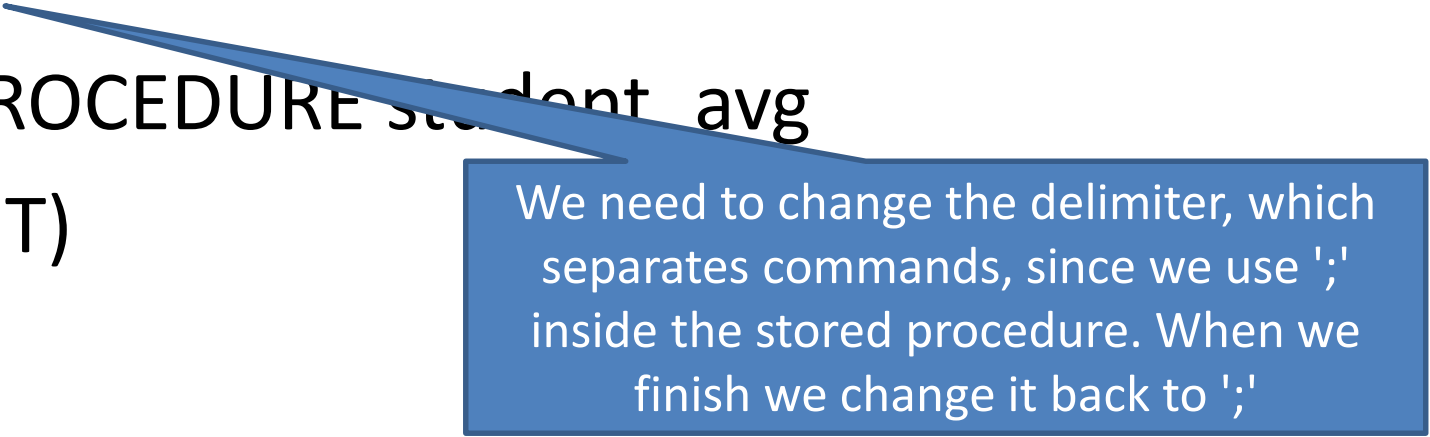
```
(IN stId INT)
```

```
BEGIN
```

```
    SELECT AVG(grade) FROM grades WHERE  
studentId = stId;
```

```
END $$
```

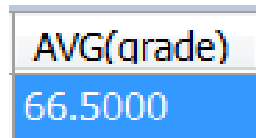
```
DELIMITER ;
```



We need to change the delimiter, which separates commands, since we use ';' inside the stored procedure. When we finish we change it back to ';'.

Stored Procedure (cont.)

- `CALL student_avg(111);`



A screenshot of a database query result. It consists of a small table with two rows. The first row has a header 'AVG(grade)' in a light gray box. The second row has the value '66.5000' in a blue box.

AVG(grade)
66.5000

- `DROP PROCEDURE student_avg`
- MySQL doesn't (really) support `ALTER PROCEDURE`, so in order to change a procedure you need to first drop it and then create it again.

Another Example

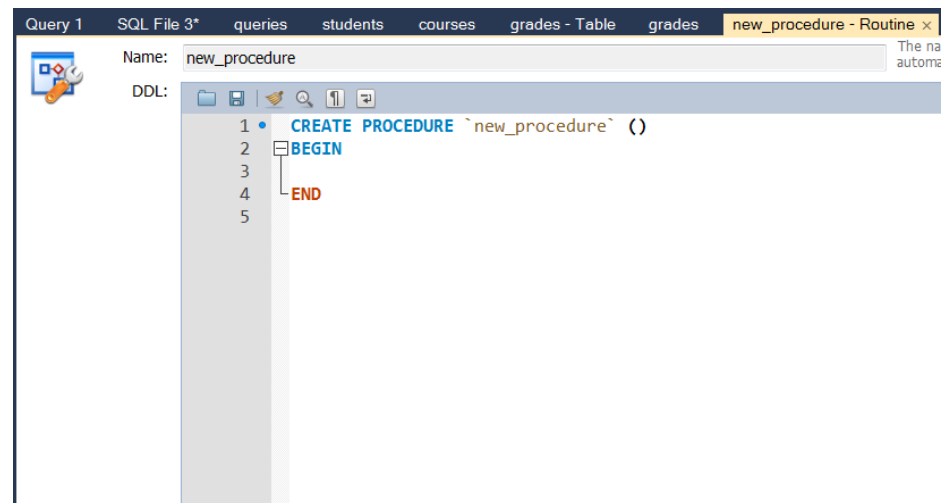
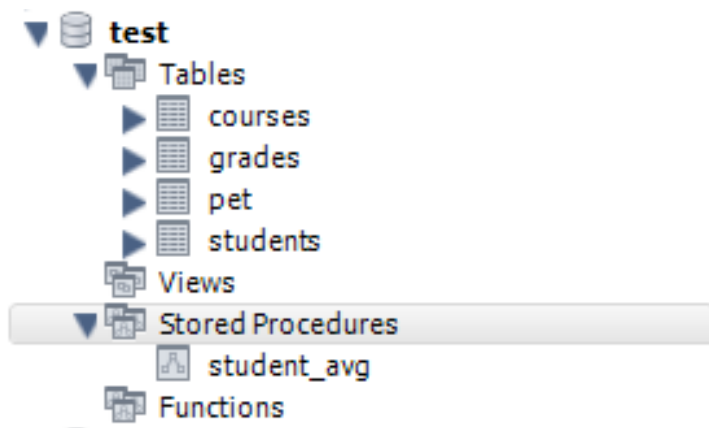
```
CREATE PROCEDURE `student_avg_2` (IN stId
INT, OUT avg_g REAL, OUT max_g INT)
BEGIN
    SELECT AVG(grade) into avg_g FROM grades
    WHERE studentId = stId;
    SELECT MAX(grade) into max_g FROM grades
    WHERE StudentId = stId;
END
```

	@avg_grade	@max_grade
	78.5	85

- CALL student_avg_2(222, @avg_grade, @max_grade);
- SELECT @avg_grade, @max_grade;

Stored Procedure Workbench

- In Workbench you can easily create and alter stored procedures using the GUI.
 - Simply right click on the "Stored Procedures" and select "Create" to create a stored procedure.
 - Right click on a stored procedure and select "alter stored procedure"



Transactions and Errors in Stored Procedures

```
CREATE DEFINER=`root`@`localhost` PROCEDURE `transfer_balance` (  
    IN sender INT,  
    IN receiver INT,  
    IN trAmount REAL)  
BEGIN  
  
    DECLARE EXIT HANDLER FOR SQLEXCEPTION, SQLWARNING  
    BEGIN  
        ROLLBACK;  
        SELECT 'Error occurred' as Message;  
    END;  
  
    START TRANSACTION;  
    SELECT @sBl = amount FROM bankBalances WHERE userId = sender;  
    UPDATE bankBalances SET amount = @sBl - trAmount WHERE userId = sender;  
    SELECT @rBl = amount FROM bankBalances WHERE userId = receiver;  
    UPDATE bankBalances SET amount = @rBl + trAmount WHERE userId = receiver;  
    COMMIT;  
  
END;
```



```

CREATE PROCEDURE [dbo].[PgRequestToPlaySP]
    -- Add the parameters for the stored procedure here
    @workerId nchar(20),
    @assignmentId nchar(40),
    @agentPlay int,
    @incriminateMode bit,
    @isIn bit out
AS
BEGIN
    SET NOCOUNT ON;
    --if has entry, update last seen
    declare @hasEntry as int;
    declare @expNum as int;
    set @expNum = 41;
    select @isIn = 1 from PgGames where workerId=@workerId and assignmentId=@assignmentId and status = 1;
    if (@isIn = 1)
    begin
        return @@Error;
    end
    set @isIn = 0;
    select @hasEntry = count(*) from PgWaitingList where workerId=@workerId and assignmentId=@assignmentId and expNum=@expNum;
    if (@hasEntry > 0)
    begin
        update PgWaitingList set lastSeen = CURRENT_TIMESTAMP where workerId=@workerId and assignmentId=@assignmentId and
expNum=@expNum;
    end
    else
    begin
        insert into PgWaitingList (expNum,workerId,assignmentId,insertTime,lastSeen) values
        (@expNum,@workerId,@assignmentId,CURRENT_TIMESTAMP,CURRENT_TIMESTAMP);
    end
    --select top 4 desc and create new games
    declare @numOfWaiting as int;
    delete from PgWaitingList where lastSeen < DateADD(mi, -1, Current TimeStamp); --remove old players
    select @numOfWaiting = count(*) from PgWaitingList where expNum=@expNum;
    if (@numOfWaiting >=4-@agentPlay)
    begin
        --build new games
        declare @maxGameId as int;
        select @maxGameId= max(gameId) from PgGames;
        if (@maxGameId is null)
        begin
            set @maxGameId = 1;
        end
        if (@maxGameId < @expNum * 1000)
        begin
            set @maxGameId = @expNum * 1000;
        end
        --do I need to lock the following two queries?
        insert into PgGames (expNum, workerId, assignmentId, gameId, status, playerId, isPirate, insertTime)
        select top(4-@agentPlay) @expNum, workerId, assignmentId, @maxGameId+1, 1, -1, 0, CURRENT_TIMESTAMP from PgWaitingList
        where expNum=@expNum order by insertTime;
        delete from PgWaitingList where assignmentId in (select assignmentId from PgGames where status=1); --remove from waiting

        update top (1) PgGames set isPirate=1 where playerId=-1 and assignmentId = (select top 1 assignmentId

```

Stored procedures may be long,
and may contain IF clauses, and
WHILE clauses

Triggers

- Suppose we would like to have a column that will hold the average for every student.
- Let's add the new column:
 - ALTER TABLE students ADD avg_grade REAL;
- Let's update all rows
 - UPDATE students s JOIN (SELECT studentId, AVG(grade) AS av from grades GROUP BY studentId) AS v ON s.id=v.studentId SET s.avg_grade=v.av;
- We still need to update it every time a grade is changed, added or deleted!

Triggers (Cont.)

DELIMITER \$\$

CREATE TRIGGER new_grade_received

AFTER INSERT ON grades

FOR EACH ROW

BEGIN

UPDATE students SET avg_grade = (SELECT AVG(grade) FROM
grades WHERE studentId=NEW.studentId) where id =
NEW.studentId;

END\$\$

DELIMITER ;

You can only define one event for each trigger (so might need multiple triggers)

You can use BEFORE if you want to access the DB before the change was made

Triggers execution


- INSERT INTO grades (courseId, studentId, grade, passed) VALUES (50, 111, 87, 1);
- SELECT * FROM students;

id	age	gender	degree	firstName	lastName	avg grade
111	21	1	1	Chaya	Glass	73.33333333
222	28	1	3	Tal	Negev	NULL
333	24	0	1	Gadi	Golan	NULL
444	23	0	1	Moti	Cohen	NULL
700	26	1	2	Maya	Levi	NULL

- DROP TRIGGER new_grade_received;

VIEWS

- Simplify complex queries
- Limit data access to specific users
- Enable computed columns



View Example

```
CREATE VIEW avg_grades_view AS  
  SELECT students.firstName, AVG(grade) AS average  
  FROM grades  
  INNER JOIN students ON grades.studentId = students.id  
  GROUP BY studentId;
```

```
SELECT * FROM avg_grades_view
```

	firstName	average
▶	Chaya	73.3333
	Tal	78.5000
	Gadi	45.0000
	Moti	95.0000

```
UPDATE avg_grades_view SET average=75 WHERE firstName  
LIKE 'Chaya';
```

Error Code: 1288. The target table grades_view of the UPDATE is not updatable

Updatable View

```
CREATE VIEW full_grades_view AS  
SELECT students.firstName, studentId, courseId, grade FROM  
grades INNER JOIN students ON grades.studentId =  
students.id;
```

```
SELECT * FROM full_grades_view;
```

```
UPDATE full_grades_view SET grade = 80 WHERE firstName LIKE  
'Chaya' AND courseId=30;
```


```
SELECT * FROM grades;
```

courseId	studentId	grade	passed
20	111	43	0
30	111	80	1
50	111	87	1
20	222	85	1
40	222	72	1
40	333	45	0
30	444	95	1

firstName	studentId	courseId	grade
Chaya	111	20	43
Chaya	111	30	90
Chaya	111	50	87
Tal	222	20	85
Tal	222	40	72
Gadi	333	40	45
Moti	444	30	95

Connecting to MySQL from Java (Connector /J)

Methods to connect to MySQL Server



SELECT * FROM students (in JAVA)

```
import java.sql.*;
```

```
public class Main{
```

```
    public static void main(String[] args){
```

```
        try{
```

Reflection

```
            Class.forName("com.mysql.jdbc.Driver");
```

```
            try(Connection con = DriverManager.getConnection("jdbc:mysql://localhost:3306/test", "user", "pwd")){
```

```
                Statement stmt = con.createStatement();
```

```
                ResultSet rs = stmt.executeQuery("SELECT * FROM students");
```

```
                int numColumns = rs.getMetaData().getColumnCount();
```

```
                while (rs.next()){
```

```
                    for (int col = 1; col <= numColumns; col++){
```

```
                        System.out.print(rs.getString(col) + " ");
```

```
                    }
```

```
                    System.out.println();
```

```
                }
```

```
            } catch (Exception ex){ex.printStackTrace();
```

```
        }
```

```
    }
```

Try with resources (java 7). No need to call con.close()

rs is initially
located before
the first row

```
111 21 1 1 Chaya Glass 73.33
222 28 1 3 Tal Negev null
333 24 0 1 Gadi Golan null
444 23 0 1 Moti Cohen null
700 26 1 2 Maya Levi null
```

```
java.lang.ClassNotFoundException: com.mysql.jdbc.Driver
```

```
    at java.net.URLClassLoader.findClass(URLClassLoader.java:381)
```

```
    at java.lang.ClassLoader.loadClass(ClassLoader.java:424)
```

```
    at sun.misc.Launcher$AppClassLoader.loadClass(Launcher.java:331)
```

```
    at java.lang.ClassLoader.loadClass(ClassLoader.java:357)
```

```
    at java.lang.Class.forName0(Native Method)
```

```
    at java.lang.Class.forName(Class.java:264)
```

```
    at ariel.databases.Main.main(Main.java:19)
```

To get column names we can
rs.getMetaData().getColumnLabels()

Jar file is missing

- Solution 1: use Gradle.
- Solution 2:
 - Goto <https://dev.mysql.com/downloads/connector/j/> download jar file.
 - Create a bin folder: copy jar file into the folder
 - Add bin to libraries (in IntelliJ: Project Structure -> Libraries -> + -> JAVA -> find bin directory)

prepareStatement

- prepareStatement allows the creating of a statement with missing parameters and filling them up later.
- May be faster and can provide some level of security (especially when part of the query are obtained from user input)

```
String query = "DELETE FROM students WHERE studentId=?"  
try (PreparedStatement pstmt = con.prepareStatement(query))  
{  
    pstmt.setString(1, userId);  
    pstmt.executeUpdate();  
}
```

executeQuery(), executeUpdate(), execute()

executeQuery()	executeUpdate()	execute()
This method is used to execute the SQL statements which retrieve some data from the database.	This method is used to execute the SQL statements which update or modify the database.	This method can be used for any kind of SQL statements.
This method returns a ResultSet object which contains the results returned by the query.	This method returns an int value which represents the number of rows affected by the query. This value will be the 0 for the statements which return nothing.	This method returns a boolean value. TRUE indicates that query returned a ResultSet object and FALSE indicates that query returned an int value or returned nothing.
This method is used to execute only select queries.	This method is used to execute only non-select queries.	This method can be used for both select and non-select queries.
Ex : SELECT	Ex : DML → INSERT, UPDATE and DELETE DDL → CREATE, ALTER	This method can be used for any type of SQL statements.

Executing Stored Procedure

```
String query = "{CALL student_avg(?)}";  
CallableStatement stmt = con.prepareCall(query);  
int studentId = 222;  
stmt.setInt(1, studentId);  
ResultSet rs = stmt.executeQuery();
```