



Spark

Amos Azaria and Merav Chkroun

Spark

- Spark allows us to compute functions that are massively parallelizable, on large amounts of data.
- Spark usually relies on the map-reduce paradigm (and the Hadoop file-system).
- Spark has a Scala, Java and Python API (with Scala being the most popular).
- We will be using PySpark which is a Python interpreter for Spark
- Installation on Linux (use VirtualBox if needed):
 - Make sure you have Java installed (and JAVA_HOME set):
 - `sudo apt-get install default-jre`
 - Make sure you have python installed (is a part of Ubuntu)
 - Download Spark from: <http://spark.apache.org/downloads.html>
 - Unzip
 - (export PYSPARK_PYTHON=python3)
 - `./bin/pyspark`
- (To install on windows see:
<http://stackoverflow.com/questions/25481325/how-to-set-up-spark-on-windows>)

Resilient Distributed Dataset (RDD)

- Spark's primary abstraction is a distributed collection of items called a Resilient Distributed Dataset (RDD).
- We can create an RDD from a file (which can be written using the Hadoop file system).
 - `text_file = sc.textFile("myDir/story.txt")`
- The text file is read line by line (i.e. is split using `"\n"`).
- We can perform transforms on an RDD and get another RDD.
 - `lines_with_hello = text_file.filter(lambda line: "hello" in line)`
- We can use `collect()` to convert an RDD to a list.
- RDD resembles Streams in Java streams.

Word Count

- The file story.txt was

<https://s3.amazonaws.com/datasets/nietzsche.txt>

```
>>> text_file = sc.textFile("myDir/story.txt")
```

```
>>> word_counts = text_file.flatMap
```

```
    .map(lambda word: word.split(' '))
```

```
    .reduceByKey(lambda word, count: count + 1)
```

```
    .collect()
```

```
>>> for word, count in word_counts.collect():
```

```
    print("the word: \"%s\" appears %s times" % (word, count))
```

Result:

...

the word: "retrograde" appears 1 time(s)

the word: "grounds" appears 4 time(s)

the word: "VOUS" appears 2 time(s)

the word: "'Flatterers" appears 1 time(s)

the word: "injustice;" appears 1 time(s)

the word: "reciprocity," appears 1 time(s)

the word: "inflicted" appears 2 time(s)

the word: "limbs." appears 1 time(s)

the word: "christened" appears 2 time(s)

the word: "majority--where" appears 1 time(s)

the word: "three-fourths" appears 1 time(s)

the word: "dish," appears 1 time(s)

the word: "73." appears 1 time(s)

the word: "ENVIRONMENT," appears 1 time(s)

the word: "'honesty;" appears 1 time(s)

...

Sort by Key

- To sort by the key we (before we call collect

```
>>> word_counts = text_file  
    .map(lambda word:  
    .reduceByKey(lam  
    .sortByKey() \  
    .collect()  
>>> for word,count in word  
    print("the word: \"
```

Result:

```
the word: "" appears 2032 time(s)  
the word: ""=Man" appears 1 time(s)  
the word: ""A" appears 2 time(s)  
the word: ""AWAY" appears 1 time(s)  
the word: ""Ah," appears 1 time(s)  
the word: ""All" appears 1 time(s)  
the word: ""And" appears 2 time(s)  
the word: ""Another" appears 1 time(s)  
the word: ""Are" appears 2 time(s)  
the word: ""BIG" appears 1 time(s)  
the word: ""BY" appears 1 time(s)  
the word: ""Bad!" appears 1 time(s)  
the word: ""Be" appears 1 time(s)  
the word: ""Better" appears 1 time(s)  
the word: ""Beyond" appears 1 time(s)
```

Sort by Value

- To sort by value we then sort by the key sort in descending order

```
>>> word_counts = (text_f  
    .map(lambda  
    .reduceByKey  
    .map(lambda  
    #False is for  
    .sortByKey(Fa  
    .map(lambda  
    .collect())
```

```
>>> for word,count in wor  
    print("the word: \
```

Result:

```
the word: "the" appears 5839 time(s)  
the word: "of" appears 4560 time(s)  
the word: "and" appears 3562 time(s)  
the word: "to" appears 2716 time(s)  
the word: "" appears 2032 time(s)  
the word: "in" appears 1995 time(s)  
the word: "a" appears 1896 time(s)  
the word: "is" appears 1857 time(s)  
the word: "that" appears 1242 time(s)  
the word: "as" appears 1172 time(s)  
the word: "it" appears 908 time(s)  
the word: "for" appears 808 time(s)  
the word: "which" appears 783 time(s)  
the word: "be" appears 740 time(s)  
the word: "with" appears 665 time(s)
```

The cache action

- When reading large files, we can use the `cache()` action to load the file into memory so the subsequent operations will be executed faster.

```
>>>text_file.cache()
```

- Note that once `cache()` is called, there is no need to call it again (this is an action, not a transformation).

Bi-Grams

- In natural language processing, we many are interested in the appearances of pairs of words (bi-grams).
- For example, for the following sentence:
"I did it you did it you did it"
We get the following bi-grams count:
[((I, did), 1), ((did, it), 3), ((it, you), 2), ((you, did), 2)]
- This is useful for part of speech tagging, feature extraction, story generation, etc.
- Tri-grams are the same with triples, and in general, we call this method n-grams.

Bi-grams (zip function)

- zip is a built-in function in Python, which receives two lists (or more) and returns a list of tuples from both lists. E.g:

```
>>>zip([1, 2, 3, 6], [10, 16, 23, 57]) =  
      [(1,10), (2, 16), (3, 23), (6, 57)]
```

- *Equivalent* to the following:

```
def zip(lst1, lst2):  
    out = []  
    for i in range(min(len(lst1), len(lst2))):  
        out.append((lst1[i], lst2[i]))  
    return out
```

Bi-grams (cont.)

```
>>>def bigram(line):  
    words = line.split()  
    return zip(words, words[1:])
```

Result:

```
[('of', 'the'), 910), ('in', 'the'), 498), ('to', 'the'),  
327), ('it', 'is'), 240), ('to', 'be'), 186), ('of', 'a'),  
171), ('and', 'the'), 157), ('for', 'the'), 149),  
('that', 'the'), 138), ('is', 'the'), 131]]
```

```
>>>pairs = text_file.flatMap(bigram)  
>>>count = pairs.map(lambda x: (x, 1)).reduceByKey(lambda a, b: a + b)  
  
>>>print(count.collect()[0:10])  
>>>print(count.map(lambda xy: (xy[1],xy[0]))  
    .sortByKey(False)  
    .map(lambda xy: (xy[1],xy[0]))  
    .collect()[0:10])
```

HDFS

- Hadoop File-System is a highly distributed file system.
- We will not be using HDFS, but for any *real* use of Spark, it is recommended to install Hadoop in order to support HDFS.

Running a Python script using Spark

- In order to import Spark into Python we need first to define the environment variable `SPARK_HOME`:
 - `export SPARK_HOME=.`
- Then we need to install findspark
 - `pip install findspark`
- Then, in the script we write:
 - `import findspark`
 - `findspark.init()`
 - `from pyspark import SparkContext`
 - `sc = SparkContext("local[*]", "Simple App")`
 - ...

* Means the maximum number of cores.
You can replace it with any other number.

Pi Estimation

```
import random
samples = 10**8
def inside(p):
    x, y = random.random(), random.random()
    return x*x + y*y < 1
hits = sc.parallelize(range(0, samples)) \
    .filter(inside).count()
print("Pi is approx. %f" % (4.0 * hits / samples))
```

Result:

Pi is approx. 3.141091

Java Streams Equivalent

```
static boolean inside(Object p)
```

```
{  
    Random rand = new Random();  
    double x = rand.nextDouble();  
    double y = rand.nextDouble();  
    return x * x + y * y < 1;  
}
```

```
public static void main(String[] args)
```

```
{  
    int samples = (int)1E+8;  
    long start_time = System.nanoTime();  
    long hits = IntStream.rangeClosed(0, samples).filter(Main::inside).count();  
    long end_time = System.nanoTime();  
    System.out.println("Pi is approx. " + (4.0 * hits / samples));  
    System.out.println("Execution took: " + ((end_time-start_time)*1E-9) + " seconds");  
}
```

Adding .parallel() did not help

Result:

Pi is approx. 3.14139924

Execution took: 8.663336708000001 seconds

When to Use Spark?

- On my laptop, using the Pi estimation example, Java Streams actually outperformed Spark, however when using a workstation (with 24 cores) Spark was twice as fast.
- Spark is intended for use with systems that:
 - Run on multiple machines / many cores
 - Have the data spread-out on multiple machines (HDFS = big data!)
 - Highly scalable

Data-frames

- Data-frames are the equivalent of Spark to tables in relational databases.
- We can create a data-frame from:
 - A relational database (and similar DBMS such as Cassandra)
 - An RDD
 - A CSV file
 - XML / JSON and other sources
- We can use SQL(!) to query Data-frames.

Data-frames example

```
from pyspark.sql import Row
from pyspark.sql import SQLContext
sqlContext = SQLContext(sc)
student_list = [(111, 'Chaya', 'Glass', 21), (222, 'Tal', 'Negev', 28),
                 (333, 'Gadi', 'Golan', 24), (444, 'Moti', 'Cohen', 23)]
student_rdd = sc.parallelize(student_list)
students_rows = student_rdd.map(lambda x:
    Row(id=int(x[0]),age=int(x[3]), firstName=x[1], lastName=x[2]))
df_students = sqlContext.createDataFrame(students_rows)
df_students.show()
```

```
+---+-----+---+-----+
|age|firstName|id|lastName|
+---+-----+---+-----+
| 21|    Chaya|111|    Glass|
| 28|     Tal|222|   Negev|
| 24|    Gadi|333|    Golan|
| 23|    Moti|444|    Cohen|
+---+-----+---+-----+
```

Create Data-frame from json

```
students_json =
```

```
'[{"id":"111","firstName":"Chaya","lastName":"Glass",  
"age":23}, {"id":"222","firstName":"Tal","lastName":"Negev",  
"age":28}, {"id":"333","firstName":"Gadi",  
"lastName":"Golan", "age":24}, {"id":"444","firstName":"Moti",  
"lastName":"Cohen", "age":23}]'
```

```
df =
```

```
sqlContext.read.json(sc.parallelize([students_json]))
```

```
df.show()
```

```
+--+-----+--+-----+  
|age|firstName|id|lastName|  
+--+-----+--+-----+  
| 21|      Chaya|111|    Glass|  
| 28|        Tal|222|    Negev|  
| 24|      Gadi|333|    Golan|  
| 23|      Moti|444|    Cohen|  
+--+-----+--+-----+
```

Create DF from json file

```
df = sqlContext.read.json("myDir\students.json",  
multiLine=True)  
df.show()
```

```
+---+-----+---+-----+  
|age|firstName|id|lastName|  
+---+-----+---+-----+  
| 21|    Chaya|111|    Glass|  
| 28|     Tal|222|   Negev|  
| 24|    Gadi|333|    Golan|  
| 23|    Moti|444|    Cohen|  
+---+-----+---+-----+
```

.withColumn()

- .withColumn() allows us to add an additional column (which is usually dependant on other columns).

```
df_students2 = df_students.withColumn('age_plus_id',  
df_students.age + df_students.id)  
df_students2.show()
```

```
+---+-----+---+-----+-----+  
|age|firstName|id|lastName|age_plus_id|  
+---+-----+---+-----+-----+  
| 21|    Chaya|111|    Glass|        132|  
| 28|     Tal|222|   Negev|        250|  
| 24|    Gadi|333|    Golan|        357|  
| 23|    Moti|444|    Cohen|        467|  
+---+-----+---+-----+-----+
```

.withColumn() (cont.)

- The second argument to withColumn must be a column, so if we want to enter our own data, it must be converted to a column.

```
from pyspark.sql import functions
from pyspark.sql.functions import lit #lit for literal
df_students2 = df_students.withColumn('young',
    functions.when(df_students.age < 25,
        True).otherwise(False)).withColumn('max_grade', lit(100))
df_students2.show()
```

```
+---+-----+---+-----+---+-----+
|age|firstName|id|lastName|young|max_grade|
+---+-----+---+-----+---+-----+
| 21|    Chaya|111|    Glass| true|      100|
| 28|     Tal|222|   Negev|false|      100|
| 24|     Gadi|333|    Golan| true|      100|
| 23|     Moti|444|    Cohen| true|      100|
+---+-----+---+-----+---+-----+
```

SQL Query in Spark

```
df_students.registerTempTable("student_table")  
ending_with_n = sqlContext.sql("SELECT firstName,  
lastName, id FROM student_table WHERE lastName  
LIKE '%n'")  
ending_with_n.show()
```

```
+-----+-----+---+  
|firstName|lastName| id|  
+-----+-----+---+  
|      Gadi|      Golan|333|  
|      Moti|      Cohen|444|  
+-----+-----+---+
```



Map-Reduce Example

- Suppose we have documents with the following structure:

```
{  
  _id: ObjectId("50a8240b927d5d8b5891743c"),  
  cust_id: "abc123",  
  ord_date: new Date("Oct 04, 2012"),  
  status: 'A',  
  amount: 25,  
  items: [ { sku: "mmm", qty: 5, price: 2.5 },  
            { sku: "nnn", qty: 5, price: 2.5 } ]  
}
```

SKU = Stock Keeping Unit
is an item identifier.

Can we do this in a
relational DB, With SQL?

We will need to assume a
different structure a
relational data-base.

- We would like to get the total **amount** paid by each **cust_id** with status 'A'.

Solving in spark with RDD

- We can solve it using map-reduce:

```
df = sqlContext.read.json("MyDir/orders.json",  
multiLine=True)
```

```
rdd = df.rdd
```

```
calc_amount = rdd.filter(lambda a: a.status == 'A') \  
    .map(lambda row: (row.cust_id, row.amount)) \  
    .reduceByKey(lambda a,b: a+b) \  
    .collect()
```

```
for cust_id, sum_amount in calc_amount:  
    print('%s : %s' %(cust_id, sum_amount))
```

279,840 orders RDD:
time: 5.2894 seconds

```
-  
abc123 : 75  
def123 : 70
```


Solving in Spark with Dataframe-1

- Assume we would have all the data in a json file.
can we obtain the query using spark?

```
df = sqlContext.read.json("myDir/orders.json",  
multiLine=True)
```

```
df.registerTempTable("orders_table")
```

```
amount = sqlContext.sql("SELECT sum(amount),  
cust_id FROM orders_table WHERE status == 'A'  
group by cust_id")
```

```
amount.show()
```

279,840 orders DF (select):
time: 0.2224 seconds

```
+-----+-----+  
|sum(amount)|cust_id|  
+-----+-----+  
|          75| abc123|  
|          70| def123|  
+-----+-----+
```

Solving in Spark with Dataframe-2

```
from pyspark.sql import functions
```

```
df.filter(df.status == 'A')\
```

```
    .groupBy(df.cust_id)\
```

```
    .agg(functions
```

```
        .sum(functions.col("amount")))\
```

```
    .show()
```

279,840 orders DF:
time: 1.30952 seconds

```
+-----+-----+
|sum(amount)|cust_id|
+-----+-----+
|          75| abc123|
|          70| def123|
+-----+-----+
```

RDD v. Data Frames

| RDD | Data Frames |
|----------------------------|---|
| “How” to do | “What” to do |
| Unstructured Data | Structured and semi-structured data |
| Less optimal and efficient | optimization and performance benefits available with DataFrames |